

Infrastructure as Code

Terraform & Azure

The Problem with "ClickOps"

You've Been Clicking Around the Portal...

What happens when you need to:

- Recreate your environment?
- Deploy to multiple regions?
- Set up dev, test, and prod environments?
- Document what you built?
- Audit history?

Manual approaches lead to:

- Inconsistency between environments
- "Configuration drift" over time
- No version history
- Human errors during deployment
- Knowledge trapped in someone's head

What is Infrastructure as Code?

*Managing and provisioning infrastructure through **code** instead of manual processes*

- Write configuration files describing your infrastructure
- Store them in version control (Git)
- Run tools to create/update infrastructure automatically
- Track changes just like application code

Benefits:

-  Repeatable and consistent
-  Version controlled
-  Documented by default
-  Testable
-  Fast to deploy
-  Easy to destroy and recreate

Infrastructure as Code Tools

Offerings

Tool	Type	Best For
Terraform	Declarative, Multi-cloud	Cloud-agnostic infrastructure
Azure ARM Templates	Declarative, Azure-only	Azure-native deployments
Bicep	Declarative, Azure-only	Modern Azure IaC
Pulumi	Imperative, Multi-cloud	Using programming languages
CloudFormation	Declarative, AWS-only	AWS infrastructure

Why Terraform?

The Industry Standard

- **Multi-cloud:** Works with Azure, AWS, GCP, and 1000+ providers
- **Declarative:** Describe what you want, not how to build it
- **Large community:** Extensive documentation and examples
- **Mature:** Battle-tested in production worldwide
- **Plan before apply:** See changes before they happen
- **State management:** Knows what currently exists

How Terraform Works

1. Write configuration (.tf files)
↓
2. `terraform init` (download providers)
↓
3. `terraform plan` (preview changes)
↓
4. `terraform apply` (create/update resources)
↓
5. State file updated (tracks reality)

Terraform File Structure

Typical project layout:

```
my-infrastructure/
├── main.tf          # Main resources
├── variables.tf     # Input variables
├── outputs.tf        # Output values
├── providers.tf      # Provider configuration
├── env.tfvars        # Variable values
└── modules/
    └── networking/
        ├── main.tf
        └── variables.tf
```

Azure Provider Setup

Connecting Terraform to Azure

The **Azure Provider** enables Terraform to manage Azure resources

- **Subscription ID:** Which Azure subscription
- **Tenant ID:** Your Azure AD tenant
- **Authentication:** How Terraform logs in

Variables and Outputs

Making Configuration Flexible

Variables = Inputs

- Parameterize your infrastructure
- Different values for dev/test/prod
- Avoid hardcoding values
- Can have defaults and validation

Outputs = Results

- Export important values
- Pass data between modules
- Display information after deployment
- Use in other tools

Variables go in, outputs come out!

Terraform Modules

Why use modules?

-  Reusability across projects
-  Consistency and standards
-  Abstraction of complexity
-  Easier testing
-  Team collaboration

Module Best Practices

Good modules are:

1. **Single-purpose:** Do one thing well
2. **Well-documented:** Clear README and variable descriptions
3. **Tested:** Verified to work correctly
4. **Versioned:** Use version tags in Git
5. **Flexible:** Configurable via variables
6. **Opinionated:** Encode best practices

Module Sources

Where to get modules:

Public Registry:

- `registry.terraform.io`
- Thousands of modules
- Official Azure modules from Microsoft
- Community contributions

Private Registry:

- Terraform Cloud
- Your company's internal modules (or one you've written!)

Git repositories:

- GitHub, GitLab, Bitbucket
- Version control with tags
- Good for custom modules

State File Management



Terraform's Memory

State file (`terraform.tfstate`) tracks:

- What resources exist
- Current configuration
- Resource metadata
- Dependencies between resources

Critical importance:

- Terraform compares desired state (your .tf files) vs actual state (state file)
- Determines what changes are needed
- **Losing state file = disaster!**

State File Challenges

Local state problems:

-  Not shared between team members
-  No locking (concurrent changes = corruption)
-  **Contains sensitive data!**
-  No versioning or backup
-  Hard to collaborate

For production: NEVER use local state!

Remote State Backends

Store state remotely for team collaboration:

Azure Storage (Recommended for Azure):

- Blob storage container
- Built-in encryption
- Access controls via Azure AD
- Automatic blob versioning
- State locking via blob lease

Other backends:

- Terraform Cloud (managed service)
- AWS S3
- Google Cloud Storage
- Consul, etcd (for advanced setups)



State Locking

Prevents you and I from destroying each others work

Without locking:

- Two people run `terraform apply` simultaneously
- State file gets corrupted
- Infrastructure becomes inconsistent

With locking:

- First person acquires lock
- Second person waits or gets error
- Safe sequential execution

Azure Storage provides automatic locking!

State Best Practices

1. **Always use remote state** for team projects
2. **Enable state locking**
3. **Enable versioning** on storage backend
4. **Encrypt state** at rest and in transit
5. **Restrict access** to state files (contain secrets!)
6. **Backup state** regularly
7. **Use separate states** for different environments

The Development Cycle

Step 1: Write

- Create or modify .tf files
- Define resources and configuration
- Use AI assistance for syntax and best practices

Step 2: Style!

- Terraform format

Step 3: Initialize

- Run `terraform init`
- Downloads providers and modules
- Configures backend

Step 4: Plan

- Run `terraform plan`
- See what will change
- Review before applying

Step 5: Apply

- Run `terraform apply`
- Confirm changes
- Terraform creates/updates resources
- State file updated

Step 6: Verify

- Check Azure portal
- Test deployed resources
- Review outputs

Useful Terraform Commands

Command	Purpose
<code>terraform init</code>	Initialize working directory
<code>terraform plan</code>	Preview changes
<code>terraform apply</code>	Create/update infrastructure
<code>terraform destroy</code>	Delete all resources
<code>terraform fmt</code>	Format code nicely
<code>terraform validate</code>	Check syntax
<code>terraform output</code>	Show output values
<code>terraform state list</code>	List resources in state

Resource Dependencies

Managing Relationships

Terraform automatically handles dependencies:

Implicit dependencies:

- Terraform detects when one resource references another
- Creates resources in the correct order
- Example: Container App needs Container App Environment

Explicit dependencies:

- Use `depends_on` when implicit isn't enough
- Forces creation order
- Useful for non-obvious dependencies

Dependency graph:

- Terraform builds internal graph
- Parallel creation when possible
- Sequential when required

Version Control for IaC

Repository structure:

```
infrastructure/
├── .gitignore
├── README.md
└── environments/
    ├── dev/
    │   ├── main.tf
    │   └── terraform.tfvars
    ├── test/
    └── prod/
└── components/
    └── container-app/
```

Ignore state and secrets
Documentation

.gitignore must include:

- `*.tfstate`
- `*.tfstate.backup`
- `.terraform/`
- `terraform.tfvars` (if contains secrets)

Importing Existing Resources

What if you already created resources manually?

Terraform import:

- Brings existing Azure resources into Terraform state
- You write the Terraform configuration
- Import command links it to existing resource
- Future changes managed by Terraform

Process:

1. Write Terraform config matching existing resource
2. Run `terraform import` with resource ID
3. Run `terraform plan` (should show no changes)
4. Now managed by Terraform!

AI can help: Generate Terraform from Azure resource properties!

Terraform Best Practices Summary

Code organization:

- Use modules for reusability
- Meaningful resource and variable names
- Consistent naming conventions
- Document with comments

State management:

- Always use remote backend for teams
- Enable state locking
- Version and backup state
- Separate states for environments

Security:

- Never commit secrets
- Use managed identities
- Encrypt state files
- Review all changes

Workflow:

- Always plan before apply
- Use version control
- Code review via pull requests
- Automate with CI/CD